

CS343: Operating System

Process Scheduling

Lect11 : 21st Aug 2023

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Process Scheduling

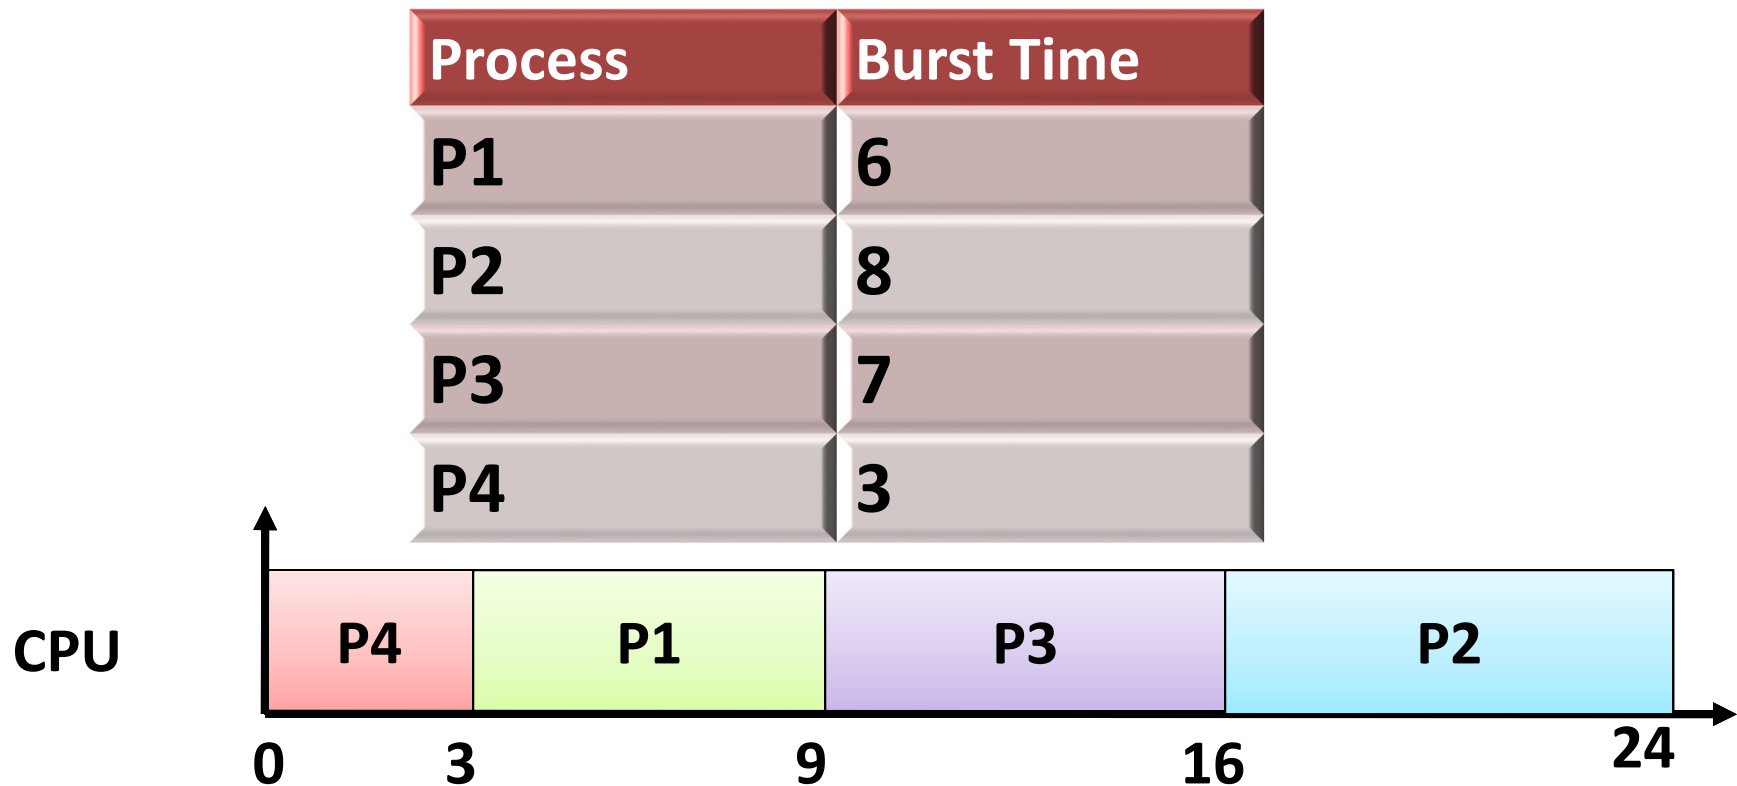
Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time

Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
- **SJF is optimal** – gives minimum average waiting time for a given set of processes
 - **The difficulty is knowing the length of the next CPU request**
 - Could ask the user

Example of SJF



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

SJF-Preemptive: Shortest-remaining-time-first


- Now we add the concepts of varying arrival times and preemption to the analysis
- **Dynamic Decision @Runtime**

| Process | Arrival Time | Burst Time |
|---------|--------------|------------|
| P1 | 0 | 8 |
| P2 | 1 | 4 |
| P3 | 2 | 9 |
| P4 | 3 | 5 |

SJF-Preemptive: Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis
- Dynamic Decision @Runtime**

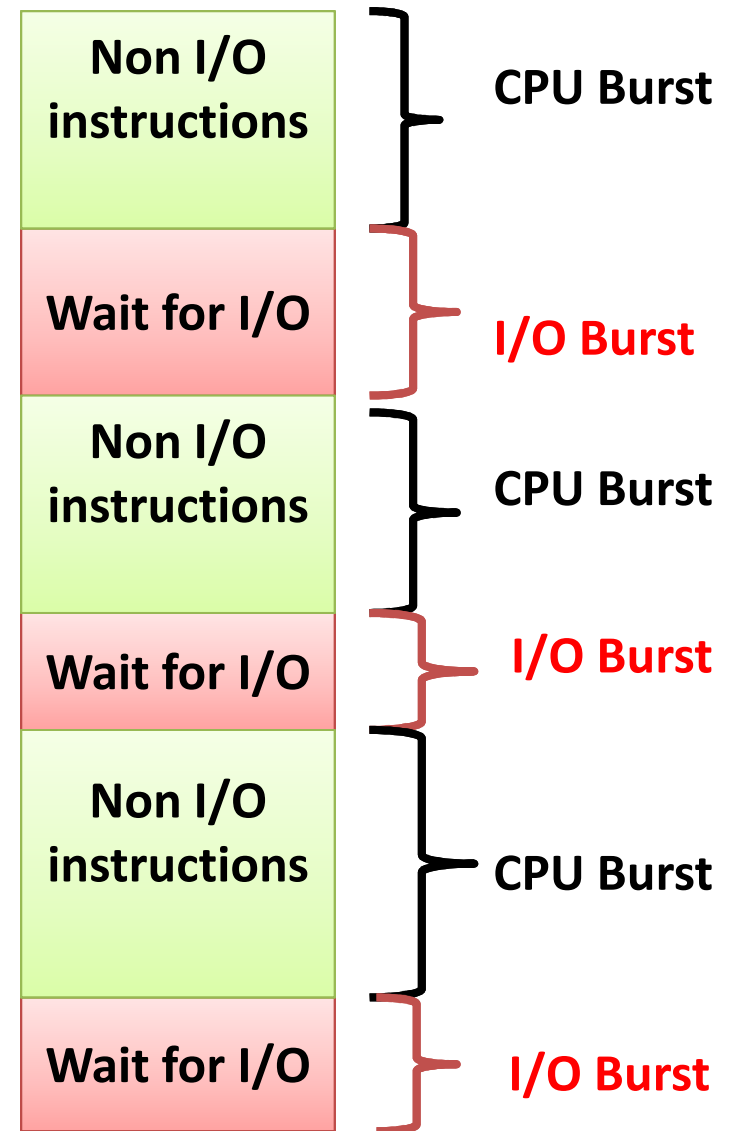
| Process | Arrival Time | Burst Time |
|---------|--------------|------------|
| P1 | 0 | 8 |
| P2 | 1 | 4 |
| P3 | 2 | 9 |
| P4 | 3 | 5 |



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3])/4 = 26/4 = 6.5$ msec

Determining Length of Next CPU Burst

- Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging



Prediction: Next CPU Burst

- Can be done by using the length of previous CPU bursts, using exponential averaging

— t_n is actual CPU Burst of nth CPU Burst

— τ_{n+1} = predicted values of next cpu burst

$$t_3 = f(t_1, t_2, t_3, \tau_1, \tau_2, \tau_3);$$

$$= \text{may be } f(t_3, \tau_3);$$

Another funny Example: http file transfer

BW depends on chunk size C_n , Initial = $1+x$,

$$C_n = (1+x)^n \text{ where } x < 1$$

EWMA Prediction: Next CPU Burst

- $\alpha, 0 \leq \alpha \leq 1$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

- Commonly, α set to $\frac{1}{2}$
- Preemptive version called **shortest-remaining-time-first**

Examples of Exponential Averaging

- $\alpha = 0$: $\tau_{n+1} = \tau_n$, Recent history does not count
- $\alpha = 1$: $\tau_{n+1} = \alpha t_n$, Only actual last CPU burst counts
- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor


Examples of Exponential Averaging

- $\alpha = 0$: $\tau_{n+1} = \tau_n$, Recent history does not count
- $\alpha = 1$: $\tau_{n+1} = \alpha t_n$, Only actual last CPU burst counts
- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$


- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$



$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n \text{ with } \alpha = 0.5$$

Prediction of the Length of the Next CPU Burst




| | | | | | | | | | |
|----------------------|----|---|---|---|----|----|----|-----|-----|
| CPU burst (t_i) | 6 | 4 | 6 | 4 | 13 | 13 | 13 | ... | |
| "guess" (τ_i) | 10 | 8 | 6 | 6 | 5 | 9 | 11 | 12 | ... |

SJF-Preemptive: Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis
- Dynamic Decision @Runtime**

| Process | Arrival Time | Burst Time |
|---------|--------------|------------|
| P1 | 0 | 8 |
| P2 | 1 | 4 |
| P3 | 2 | 9 |
| P4 | 3 | 5 |



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3])/4 = 26/4 = 6.5$ msec

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority
 - Smallest integer == > highest priority
 - Preemptive, Non-preemptive
- SJF is priority scheduling
 - Where priority is $1/\text{CPU Burst time}$

Starvation: Priority Scheduling

- **Starvation Problem**

- low priority processes may never execute
- Starve to execute...

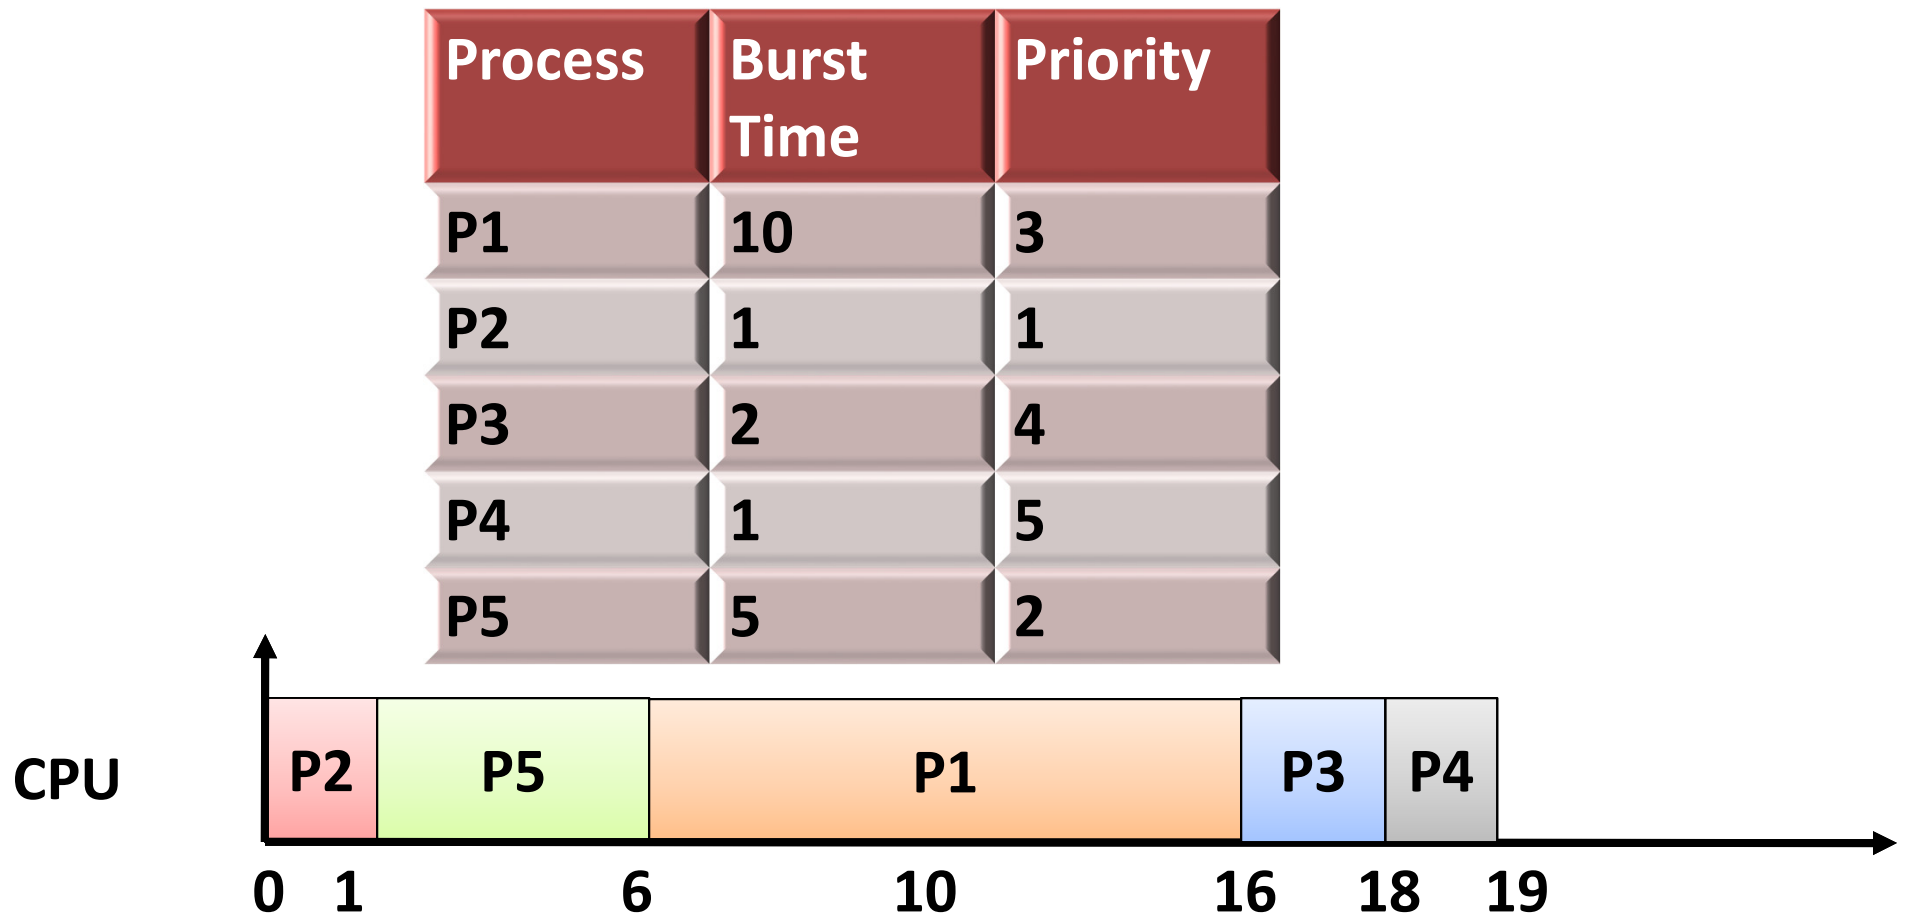
- **Solution**

- **Aging** – as time progresses increase the priority of the process

Example of Priority Scheduling

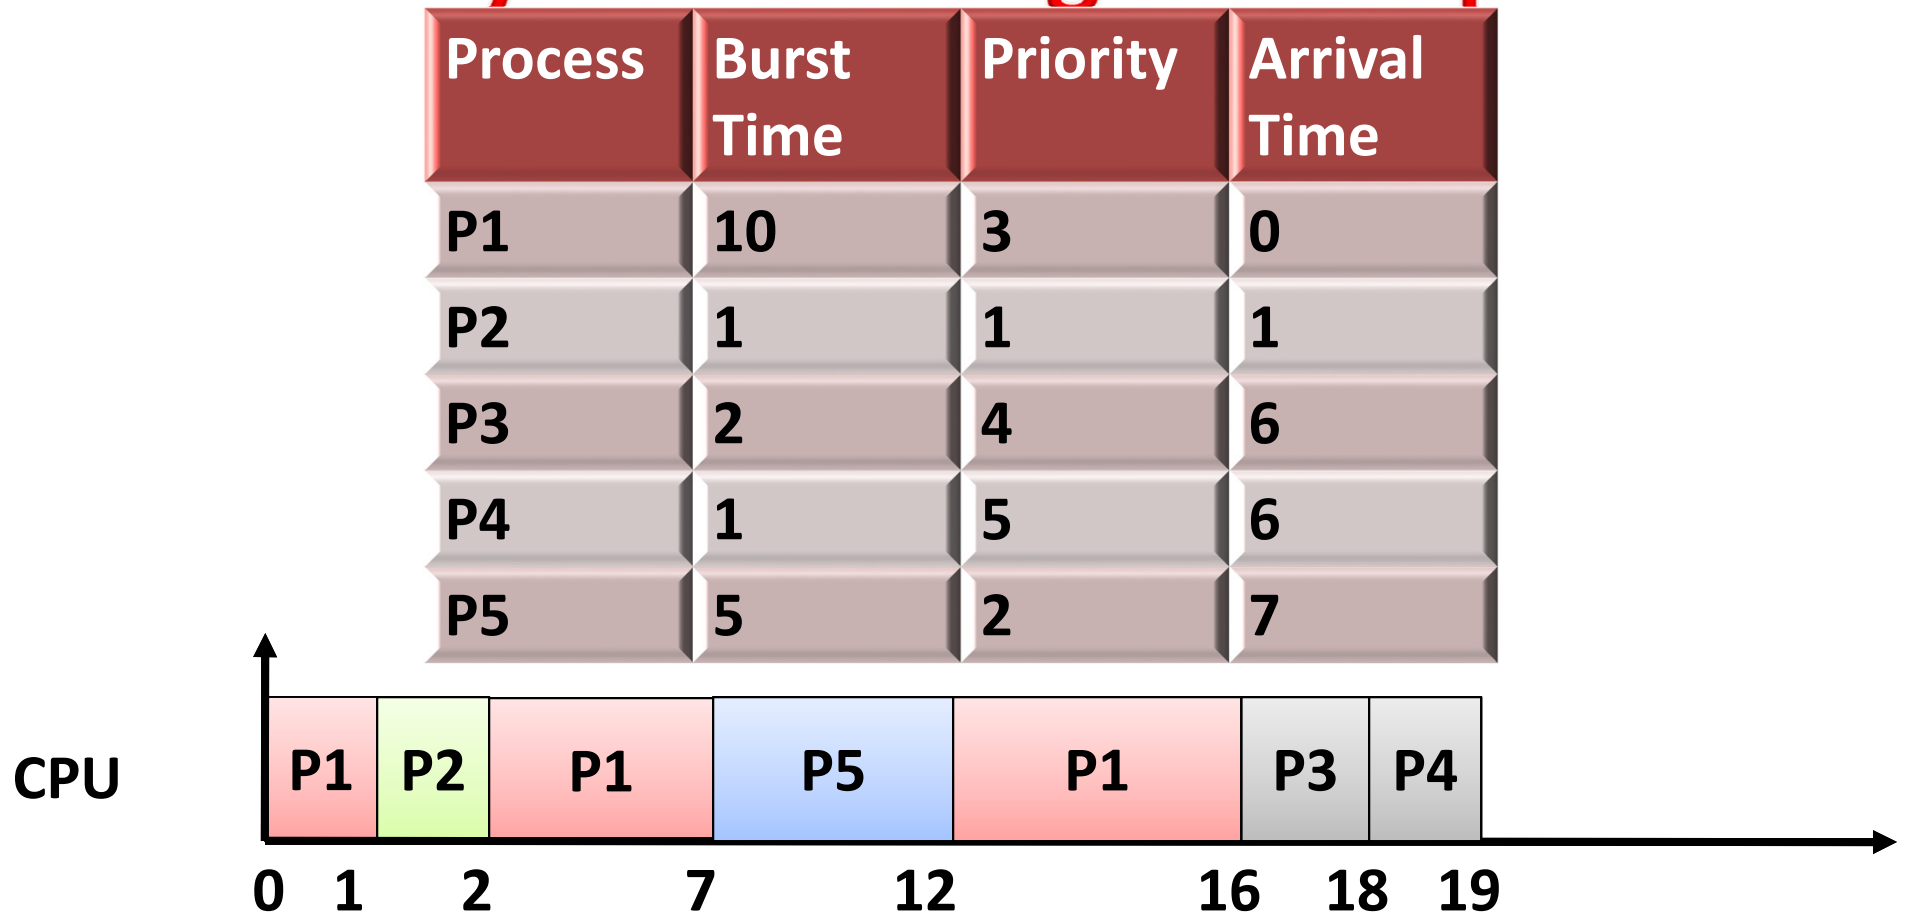
| Process | Burst Time | Priority |
|---------|------------|----------|
| P1 | 10 | 3 |
| P2 | 1 | 1 |
| P3 | 2 | 4 |
| P4 | 1 | 5 |
| P5 | 5 | 2 |

Example of Priority Scheduling



- Assume all arrived@ time 0
- Average waiting time $(6+0+16+18+1=41)/5 = 8.2$ msec

Priority Scheduling Preemptive



- All arrived@ different time
- Average waiting time $(6+0+16+18+7=47)/5$
=9.4 msec

Round Robin (RR)

- **Gol Gappe Wala Scheduling**
- Each process gets a small unit of CPU time (**time quantum q**)
 - Usually 10-100 milliseconds.
 - After this time has elapsed, the process is preempted and added to the end of the ready queue.


Round Robin (RR)

- If there are n processes in the ready queue and the time quantum is q , then
 - Each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
- No process waits more than $(n-1)q$ time units.

Round Robin (RR)

- Timer interrupts every quantum to schedule next process
 - **Timer: Hardware unit, similar to Alarm**
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high


Example : RR (with $q = 4$)



- Typically, higher Average turnaround than SJF, but better **response**
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time

Process time 10



| Quantum | Context Switches |
|---------|------------------|
| 12 | 0 |
| 6 | 1 |
| 1 | 9 |

Thanks