

CS343: Operating System

Threading and Synchronization

Lect16 : 04th Sept 2023

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Threading
- Threading Examples
- Thread mappings
 - Pthread/Uthread, Kthread, Hthread
- Synchronization

Pthread, C++ Thread, Cilk and OpenMP

Pthread

```
pthread_t tid1, tid2;  
pthread_create(&tid1, NULL, Fun1, NULL);  
pthread_create(&tid2, NULL, Fun2, NULL);  
pthread_join(tid1, NULL);  
pthread_join(tid2, NULL);
```

```
thread t1(Fun1);  
thread t1(Fun2, 0, 1, 2);  
    // 0, 1, 2 param to Fun2  
t1.join();  
t2.join();
```

C++ thread

Posix Threads (Pthreads) Interface

- **Creating and reaping threads**
 - `pthread_create`, `pthread_join`
- **Determining your thread ID** : `pthread_self`
- **Terminating threads**
 - `pthread_cancel`, `pthread_exit`
 - `exit` [terminates all threads], `return` [terminates current thread]
- **Synchronizing access to shared variables**
 - `pthread_mutex_init`,
`pthread_mutex_[un]lock`
 - `pthread_cond_init`,
`pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```
/* thread routine */  
void *HelloW(void *vargp) {  
    printf("Hello, world!\n");  
    return NULL;  
}
```

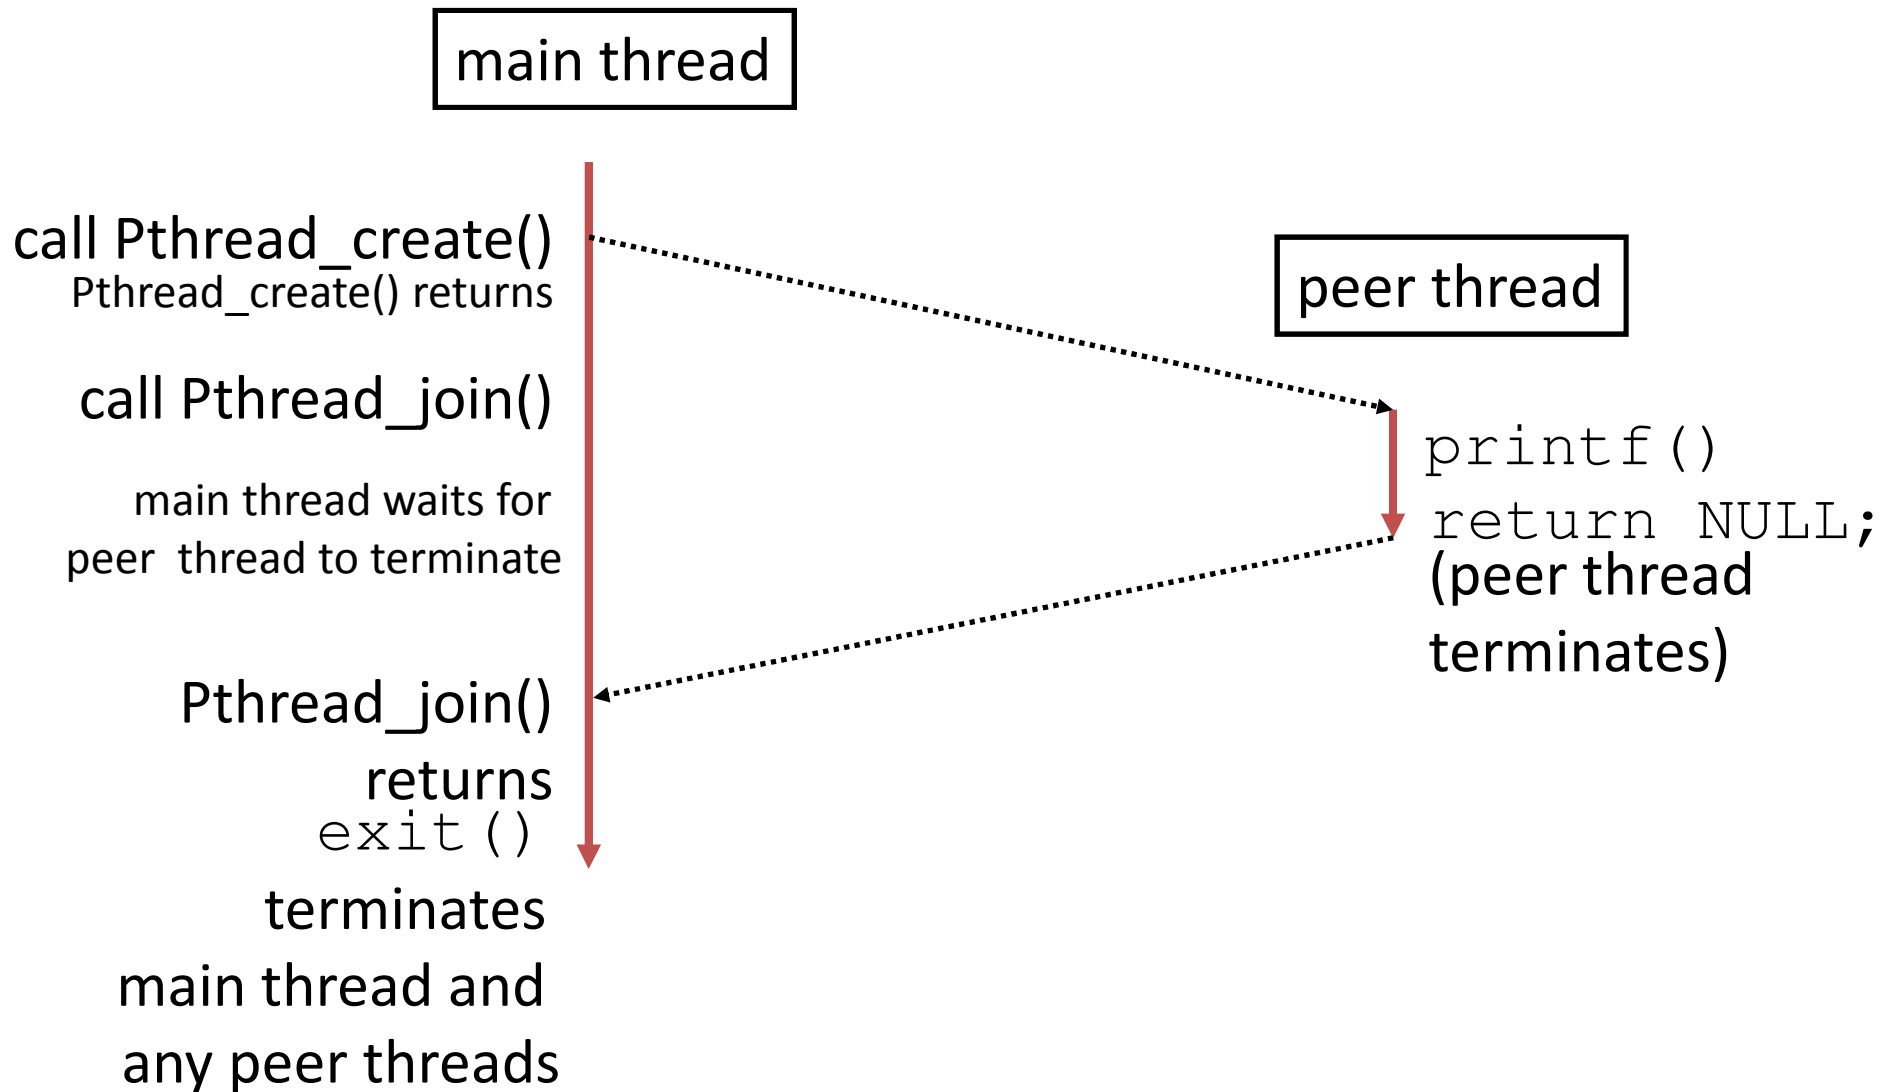
```
int main() {  
    pthread_t tid;  
    pthread_create(&tid, NULL, HelloW, NULL);  
    pthread_join(tid, NULL);  
    return 0;  
}
```

Thread attributes
(usually NULL)

Thread arguments
(void *p)

return value
(void **p)

Execution of Threaded “hello, world”



Pros and Cons: Thread-Based Designs

- **+ Easy to share data structures between threads**
 - E.g., logging information, file cache
- **+ Threads are more efficient than processes**
- **– Unintentional sharing can introduce subtle and hard-to-reproduce errors!**
 - Ease of data sharing is greatest strength of threads
 - Also greatest weakness!

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

Suppose Size=1000

0-249	250-499	500-749	750-999
--------------	----------------	----------------	----------------

T1

T2

T3

T4

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

- Independent
- Divide work into equal for each thread
- Work per thread: $\text{Size}/\text{numThread}$

VectorSum Parallel

```
void *DoVectorSum(void *tid) {  
    int j, SzPerthrd, LB, UB, TID;  
    TID= *((int *)tid);  
    SzPerthrd= (VSize/NUM_THREADS);  
    LB= SzPerthrd*TID; UB=LB+SzPerthrd;  
  
    for (j=LB; j<UB; j++)  
        A[j]=B[j]+C[j];  
}
```



VectorSum Parallel

```
int main() {  
    int i;  
    pthread_t thread[NUM_THREADS];  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_create(&thread[i],  
            NULL, DoVectorSum, (void*)&i);  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_join(thread[i], NULL);  
  
    return 0;  
}
```

Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$



Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$

- Divide the whole rows to T chunks
 - Each chunk contains : N/T rows, Assume $N\%T=0$

Matrix multiply Serial

```
void MatMul () {  
    int i, j, k, S;  
    for (i=0; i<Size; i++)  
        for (j=0; j<Size; j++) {  
            S=0;  
            for (k=0; k<Size; k++)  
                S=S+A[i][k]*B[k][j];  
            C[i][j]=S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
void * DoMatMulThread(void *arg) {  
    int i, j, k, S, LB, UB, TID, ThrdSz;  
    TID = * ((int *) arg); ThrdSz = Size / NumThrd;  
    LB = TID * ThrdSz; UB = LB + ThrdSz;  
  
    for (i = LB; i < UB; i++)  
        for (j = 0; j < Size; j++) {  
            S = 0;  
            for (k = 0; k < Size; k++)  
                S = S + A[i][k] * B[k][j];  
            C[i][j] = S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
int main() {  
    pthread_t  thread[NumThread];  
    int t;  
    Initialize();  
    for (t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
            DoMatMulThread, &t);  
    for (t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
  
    TestResult();  
    return 0;  
}
```


Estimating π using Monte Carlo


- The probability of a random point lying inside the unit circle:

$$P(x^2 + y^2 < 1) = \frac{A_{\text{circle}}}{A_{\text{square}}} = \frac{\pi}{4}$$

- If pick a random point N times and M of those times the point lies inside the unit circle:

$$P^{\circ}(x^2 + y^2 < 1) = \frac{M}{N}$$

- If N becomes very large, $P = P^{\circ}$



$$\pi = \frac{4 \cdot M}{N}$$

Value of PI: Monte-Carlo Method

```
void MontePI () {  
    int count=0,i;  
    double x,y,z;  
    for ( i=0; i<niter; i++) {  
        x = (double)rand() /RAND_MAX;  
        y = (double)rand() /RAND_MAX;  
        z = x*x+y*y;  
        if (z<=1) count++;  
    }  
    pi=(double) count/niter*4;  
}
```

PI- Multi-threaded

- 1 thread you are able to generate N points
 - Suppose M points fall under unit circle
 - $PI = 4M/N$
- With 10 thread generate 10XN points and calculate more accurately
 - Each thread calculate own value of PI (or M)
 - Average later on (or recalculate PI from collective M)

Value of PI: Pthreaded

```
int main() {  
    pthread_t  thread[NumThread]; double pi;  
    int t, at[NumThread], count, TotalIter;  
    for(t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
                        DoLocalMC_PI, &t);  
    for(t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
    for(t=0; t<NumThread; t++) count+=LCount[t];  
    TotalIter=niter*NumThread;  
    pi=((double) count/TotalIter) *4;  
    return 0;  
}
```


Value of PI: Pthreaded

```
int LCount [NumThread];  
void *DoLocalMC_Pi (void *aTid) {  
    int tid, count, i; double x, y, z;  
    tid= *((int *) aTid);  
    count=0; LCount[tid]=0;  
    for ( i=0; i<niter; i++) {  
        x = (double) rand() / RAND_MAX;  
        y = (double) rand() / RAND_MAX;  
        z = x*x+y*y; if (z<=1) count++;  
    }  
    LCount[tid]=count;  
}
```

Performance of Parallel Program (Amdahl's Law)

Example: OpenMP Parallel Program

```
printf("begin\n");  
N = 1000;  
  
#pragma omp parallel for  
for (i=0; i<N; i++)  
    A[i] = B[i] + C[i];  
  
M = 500;  
  
#pragma omp parallel for  
for (j=0; j<M; j++)  
    p[j] = q[j] - r[j];  
  
printf("done\n");
```



Speed up and efficiency


- **Notion** : T_1 = Time on Uni-processor, T_p = Time on p Processors

$$\text{Speed up} = S_p = T_1 / T_p \leq p$$


$$\text{Efficiency} = E_p = T_1 / (p \cdot T_p)$$

- Usually $S_p < p$ or $E_p < 1$ due to overhead
- Some time superliner speed up reported ($S_p > p$ or $E_p > 1$)
 - Failure to use the best sequential algorithm
 - Advantage due to larger memory


Amdahl's law




Work 500,
Time 500
Sp=1X



Work 500,
Time 400
Sp=1.25X



Work 500,
Time 350
Sp=1.4X



Work 500,
Time 300
Sp=1.7X

Amdahl's Law

$$\text{Serial fraction} = s = \frac{T_s}{T_1}$$

$$T_p = T_s + \frac{T_1 - T_s}{p}$$

$$S_p = \frac{T_1}{T_p} = \frac{T_1}{T_s + \frac{T_1 - T_s}{p}} = \frac{T_1}{T_s(1 - \frac{1}{p}) + \frac{T_1}{p}}$$

Amdahl's Law

$$S_p = \frac{T_1}{T_s(1 - \frac{1}{p}) + \frac{T_1}{p}}$$


$$Sp = \frac{1}{s(1 - \frac{1}{p}) + \frac{1}{p}} = \frac{p}{s(p-1) + 1}$$

$$s = \frac{T_s}{T_1}$$

Lt

$p \rightarrow \infty$

$$S_p = \frac{1}{s}$$



Assumption behind Amdahl's Law

- All the processors are homogeneous
- All the communication costs are zero
- All the memory accesses takes unit time (PRAM)
- All the parallel section are purely parallel:
Divisible load

$$S_p = \frac{1}{s \left(1 - \frac{1}{p}\right) + \frac{1}{p}} =$$

$$p \rightarrow \infty \quad S_p = \frac{1}{s}$$

User Threads and Kernel Threads

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**, Windows threads, Java threads
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
 - Windows , Solaris, Linux

User Threads and Kernel Threads

Kernel Thread

- Kernels are generally multithreaded (kthread)
- Pthread : user level thread
- To make concurrency cheaper the execution aspect of process is separated out into threads.
- OS manages/schedule threads and processes
- All thread operations are implemented in the kernel
- OS managed threads are called kernel-level threads or light weight processes
 - Window NT: Thread, Solaris: LWP

Kernel Thread

- The kernel knows about and manages the threads
- No runtime system is needed in this case.
- Kernel
 - A thread table that keeps track of all threads in the system.
 - In addition, process table to keep track of processes.
- OS kernel provides system call to create and manage threads

Advantage of Kernel Thread

- **ADV: As kernel has full knowledge of all threads**
 - Scheduler may decide to give more time to a process having large number of threads than process having small number of threads.
- **ADV: kthreads are especially good for apps that frequently block.**

DisAdv of Kernel Thread

- **DISADV: The kthreads are slow and inefficient.**
 - threads operations are hundreds of times slower than that of user-level threads.
- **DISADV: Since kernel must manage and schedule threads as well as processes.**
 - It require a full thread control block (TCB) for each thread to maintain information about threads.
 - As a result there is significant overhead and increased in kernel complexity.

User level thread

- Kthreads make concurrency much cheaper than process
 - because, much less state to allocate and initialize.
- However, for fine-grained concurrency kthreads still suffer from too much overhead.
 - Thread operations still require system calls.
 - Ideally, we require thread operations to be as fast as a procedure call.
- For fine grained concurrency we need "cheaper" threads.
- To make threads cheap and fast, they need to be **implemented at user level.**

User level thread

- User-Level threads are managed
 - Entirely by the run-time system (user-level library).
- Kernel knows nothing about user-level threads
 - Manages them as if they were single-threaded processes.
- User-Level threads are small and fast

User level thread : pthread

- Each **Uthread** is represented by a PC, register, stack, and small thread control block.
- Done via procedure call. i.e **no kernel involvement**
 - Creating a new thread, switching between threads, and synchronizing threads are
- Uthreads are **100x** faster than Kthreads.

Uthread : Advantage

- **The most obvious advantage of this technique**
 - uthread package can be implemented on an OS that does not support kthreads.
 - uthreads does not require modification to OS
- **Simple Representation**
 - Each thread is represented simply by a PC, registers, stack and a small control block TCB
 - all stored in the user process address space.
- **Simple Management**
 - Creating a thread, thread switching and synch threads
 - All be done without intervention of the kernel.
- **Fast and Efficient**
 - Thread switching is not much more expensive than a procedure call

Uthread: Disadvantage

- Uthreads are not a perfect solution as with everything else, they are a trade off.
 - Since, Uthreads are invisible to the OS they are not well integrated with the OS, As a result, Os can make poor decisions
 - Scheduling a process with idle threads, blocking a process whose thread initiated an I/O even though the process has other threads that can run and unscheduling a process with a thread holding a lock.
 - Solving this requires communication between between kernel and user-level thread manager.

Uthread: Disadvantage

- Lack of coordination between threads and OS
 - Process as whole gets one time slice irrespective of whether process has one thread or 1000 threads within.
- Uthreads requires ***non-blocking systems call*** i.e., a multithreaded kernel.
 - Otherwise, entire process will be blocked in the kernel, even if there are runnable threads left in the processes.
 - For example, if one thread causes a page fault, the process blocks.