

CS343: Operating System

Threading and Synchronization

Lect18 : 08th Sept 2023

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Threading
- Threading Examples
- Thread mappings
 - Pthread/Uthread, Kthread, Hthread
- Synchronization

Posix Threads (Pthreads) Interface

- **Creating and reaping threads**
 - `pthread_create`, `pthread_join`
- **Determining your thread ID** : `pthread_self`
- **Terminating threads**
 - `pthread_cancel`, `pthread_exit`
 - `exit` [terminates all threads], `return` [terminates current thread]
- **Synchronizing access to shared variables**
 - `pthread_mutex_init`,
`pthread_mutex_[un]lock`
 - `pthread_cond_init`,
`pthread_cond_[timed]wait`

User Threads and Kernel Threads

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**, Windows threads, Java threads
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
 - Windows , Solaris, Linux

User Level Thread

- **Advantages**

- User level threads are simpler and faster to generate. They are also easier to manage.
- Thread switching in user-level threads doesn't need kernel mode privileges.
- These are more portable.
- These threads may be run on any OS.

- **Disadvantages**

- The complete process is blocked if a user-level thread runs a blocking operation.
- User-level threads don't support system-wide scheduling priorities.
- It is not appropriate for a multiprocessor system.

Kernel Level Thread

- **Advantages**

- If a thread in the kernel is blocked, it does not block all other threads in the same process.
- Several threads of the same process might be scheduled on different CPUs in kernel-level threading.
- Kernel routines can be multithreaded as well.

- **Disadvantages**


- Compared to user-level threads, kernel-level threads take longer to create and maintain.
- A mode switch to kernel mode is important to transfer control from one thread in a process to another.

Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many


Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - **Solaris Green Threads**
 - **GNU Portable Threads**




One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples: Windows, Linux, Solaris 9 and later




Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows with the *ThreadFiber* package



Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX, HP-UX, Tru64 UNIX
 - Solaris 8 and earlier



Multithreaded program on Multi-core

- Uthread, Kthread
- Hthread (hardware thread) : **MIT term HART**
 - **Provided by hardware CPU**
- Mapping Kthread to Hthread (Same model)
- Mapping Uthread to Hthread (By passing Kthread) HART
 - *Allow creation an Uthread if a Hthread is free*
- **Thread affinity**
 - Suppose a thread is suitable to a particular Hthread
 - Benefit of cache and resources

Running a Apps on Specific core

```
#define _GNU_SOURCE
#include <pthread.h> #include <stdio.h> #include <stdlib.h> #include <errno.h>
#define handle_error_en(en, msg) \
    do { errno = en; perror(msg); exit(EXIT_FAILURE); } while (0)
int MYBenchmark() { while(1); }

int main(int argc, char *argv[]){ // a.out 3 # run the apps on core id =3
    int Proc,s; // Check on System Monitor Apps
    cpu_set_t cpuset;
    pthread_t thread;
    thread = pthread_self();
    Proc=atoi(argv[1]);
    CPU_ZERO(&cpuset);
    CPU_SET(Proc, &cpuset);
    s = pthread_setaffinity_np(thread, sizeof(cpu_set_t), &cpuset);
    MYBenchmark();
    exit(EXIT_SUCCESS);
}
```

Shared Variable

```
#define NTH 10
```

```
int counter=0;//Shared among thread
```

```
void SimpleCnt () {
```

```
    for(int i=0;i<10;i++)
```

```
        counter++;
```

```
}
```

Shared Variable: SimpleCnt

```
int main() {  
    thread_t th[NTH];  
    int i, j;  
    for(i=0; i < NTH; i++) {  
  
        thread_create(&th[i], NULL, SimpleCnt, NULL  
        );  
    }  
    for(j=0; j < NTH; j++) {  
        pthread_join( th[j], NULL);  
    }  
    printf("Final Ctr val: %d\n", counter);  
}
```


Shared Variable with Mutex

```
#define NTH 10  
pthread_mutex_t M = 0;  
void MutexCnt () {  
    for (int i=0; i<10; i++) {  
        pthread_mutex_lock( &M );  
        counter++;  
        pthread_mutex_unlock( &M );  
    }  
}
```


Shared Variable: MutexCnt

```
int counter = 0;
int main() {
    thread_t th[NTH];
    int i, j;
    for(i=0; i < NTH; i++) {
        thread_create(&th[i], NULL, MutexCnt, NULL);
    }
    for(j=0; j < NTH; j++) {
        pthread_join(th[j], NULL);
    }
    printf("Final Ctr val: %d\n", counter);
}
```

Many threads trying to acquire Mutex LOCK



Shared Variable with Mutex

```
#define NTH 10

pthread_mutex_t M = 0;

void MutexCnt() {
    int local_cnt=0
    for(int i=0;i<10;i++) local_cnt++;
    pthread_mutex_lock( &M );
    counter +=local_count;
    pthread_mutex_unlock( &M );
}
```

Synch. Primitives

- `pthread_mutex_init, lock, unlock, trylock`
- `pthread_attr_set detachstate, guardsize_np, stacksize, inheritsched, schedpolicy, schedparam`
- `pthread_cond_wait, signal, broadcast, init, destroy`
- Our main concern
 - Lock, unlock, trylock, condsignal, condbroadcast

Locking Overhead

- Serialization points
 - Minimize the size of critical sections
 - Be careful
- Rather than wait, check if lock is available
 - **pthread_mutex_trylock**
 - If already locked, will return EBUSY
 - Will require restructuring of code
 - **Suspend self by pthread_yield() Give chance to others**
 - **Suspend self by doing a timed wait..**

1. Process and Thread Interchangeably
Processes/Threads share data

2. Mutex Via Semaphore

Semaphore

- Earlier : Shared thread incrementing counter
- Train : One track many trains
 - **MUT**ual **EX**clusion is must
 - **Train Collision**

Life loss



Semaphore

- Semaphore : A system of sending messages by holding The arms or two flags or poles in certain positions according to an alphabetic code




Semaphore

- Semaphore : A system of sending messages by holding The arms or two flags or poles in certain positions according to an alphabetic code



Semaphore

- Semaphore : A system of sending messages by holding The arms or two flags or poles in certain positions according to an alphabetic code



Locking Overhead

- Serialization points
 - Minimize the size of critical sections
 - Be careful
- Rather than wait, check if lock is available
 - **pthread_mutex_trylock**
 - If already locked, will return EBUSY
 - Will require restructuring of code
 - **Suspend self by pthread_yield() Give chance to others**
 - **Suspend self by doing a timed wait..**


Synchronization Hardware and Algorithms

Data Consistency

- Processes/thread can execute concurrently
 - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Classic Example: Producer & Consumer

- There is a buffer of size `BUFFER_SIZE`
- When a producer : produce an Item, he put into Buffer and increment the counter
- When a consumer : consume an Item, he read an Item from Buffer and decrement the counter
- Initially, **counter** is set to 0



Producer

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```


Assumption

- Assume that all the instruction are atomic
 - The **load** and **store** machine-language instructions are atomic;
 - That is, cannot be interrupted
- We will see: Still we have problem in Synchronization
 - Or some protocol/Algorithm to Handle
 - We may need different hardware support, **a specific kind of Instruction to be atomic**

Race Condition: ProducerConsumer

Counter ++

could be implemented as

```
reg1 = counter //LD  
reg1 = reg1 + 1 //INC  
counter = reg1 //ST
```

Counter - -

could be implemented as

```
reg2 = counter //LD  
reg2 = reg2 - 1 //DEC  
counter = reg2 //ST
```

Consider this execution interleaving with “count = 5” initially:

S0: producer execute	reg1 = counter	{reg1 = 5}
S1: producer execute	reg1 = reg1 + 1	{reg1 = 6}
S2: consumer execute	reg2 = counter	{reg2 = 5}
S3: consumer execute	reg2 = reg2 - 1	{reg2 = 4}
S4: producer execute	counter = reg1	{counter = 6 }
S5: consumer execute	counter = reg2	{counter = 4}

Expected Result =5

Critical Section Problem

- Consider system of n processes

$$\{p_0, p_1, \dots, p_{n-1}\}$$


- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc
 - When one process in critical section, no other may be in its critical section

Critical Section Problem

- ***Critical section problem*** is to design protocol to solve this
 - Each process must ask permission to enter critical section in **entry section**,
 - May follow critical section with **exit section**, then **remainder section**

Critical Section

- General structure of process P_i



Thanks