

CS343: Operating System

Synchronization

Lect19 : 12th Sept 2023

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Synchronization
 - Critical Section Problem
- Solution to CS Problems
 - Two Threads Solutions: Peterson's Solution
 - N Thread Solutions: Filter and Bakery Algorithms
- Sync Hardware
 - CAS, TAS, LL-LC, BackupLock

Critical Section Problem

- Consider system of n processes

$$\{p_0, p_1, \dots, p_{n-1}\}$$


- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc
 - When one process in critical section, no other may be in its critical section

Critical Section Problem

- ***Critical section problem*** is to design protocol to solve this
 - Each process must ask permission to enter critical section in **entry section**,
 - May follow critical section with **exit section**, then **remainder section**

Critical Section

- General structure of process P_i



Solution to Critical-Section Problem

- **Mutual Exclusion**
- **Progress**
- **Bounded Waiting**

Solution to Critical-Section Problem

- **Mutual Exclusion**

- If process P_i is executing in its CS
- Then no other processes can be executing in their CS

- **Progress : Deadlock free**

- If no process is running in its CS and there exist some processes that wish to enter their CS,
- Then the selection of the processes that will enter the CS next cannot be postponed indefinitely

Solution to Critical-Section Problem

- **Bounded Waiting : Starvation**

- A bound must exist on the number of times that other processes are allowed to enter their CSs after a process has made a request to enter its CS and before that request is granted
- Assume that each process executes at a nonzero speed
- No assumption concerning **relative speed** of the n processes

Critical-Section Handling in OS

- Two approaches depending on if kernel is preemptive or non- preemptive
 - **Preemptive** – allows preemption of process when running in kernel mode
 - **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
 - Essentially free of race conditions in kernel mode

Locking Algorithms

Ref: Galvin Book and
*“Art of Multiprocessor
Programming”* by Herlihy and
Shavit


Events

- An *event* a_0 of thread A is
 - Instantaneous
 - No simultaneous events (break ties)





Intervals

- An *interval* $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1



Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either

  or  

$CS_i^k \rightarrow CS_j^m$

$CS_j^m \rightarrow CS_i^k$

Two-Thread vs n -Thread Solutions

- 2-thread solutions first
 - Illustrate most basic ideas
 - Fits on one slide
- Then n -thread solutions

LockOne : Algorithm 1

```
class LockOne {  
    bool flag[2]={0,0};  
    public void lock() {  
        int j,i=ThreadID.get(); j=1-i;  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        int i = ThreadID.get();  
        flag[i] = false;  
    }  
}
```

LockOne : Algorithm 1

```
class LockOne {  
    bool flag[2]={0,0};  
    public void lock() {  
        int j,i=ThreadID.get();j=1-i;  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        int i = ThreadID.get();  
        flag[i] = false;  
    }  
}
```

Each thread
has flag

LockOne : Algorithm 1

```
class LockOne {  
    bool flag[2]={0,0};  
    public void lock() {  
        int j,i=ThreadID.get(); j=1-i;  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        int i = ThreadID.get();  
        flag[i] = false;  
    }  
}
```

Set my flag

LockOne : Algorithm 1

```
class LockOne {  
    bool flag[2]={0,0};  
    public void lock() {  
        int j,i=ThreadID.get(); j=1-i;  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        int i = ThreadID.get();  
        flag[i] = false;  
    }  
}
```

Wait for other flag to
become **false**

LockOne : Algorithm 1

```
class LockOne {  
    bool flag[2]={0,0};  
    public void lock() {  
        int j,i=ThreadID.get(); j=1-i;  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        int i = ThreadID.get();  
        flag[i] = false;  
    }  
}
```

I have released

LockOne Satisfies Mutual Exclusion

- Assume CS_A^j overlaps CS_B^k
- Consider each thread's last (j-th and k-th) read and write in the lock() method before entering
- Derive a contradiction

From the Code

- $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow$
 $\text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \text{CS}_A$
- $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
 $\text{read}_B(\text{flag}[A]==\text{false}) \rightarrow \text{CS}_B$

```
public void lock() {  
    flag[i] = true;  
    while (flag[j]) {}  
}
```

From the Assumption

- **Load Store are atomic**
 - **It will behaves the Partial Order**
- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

Combining

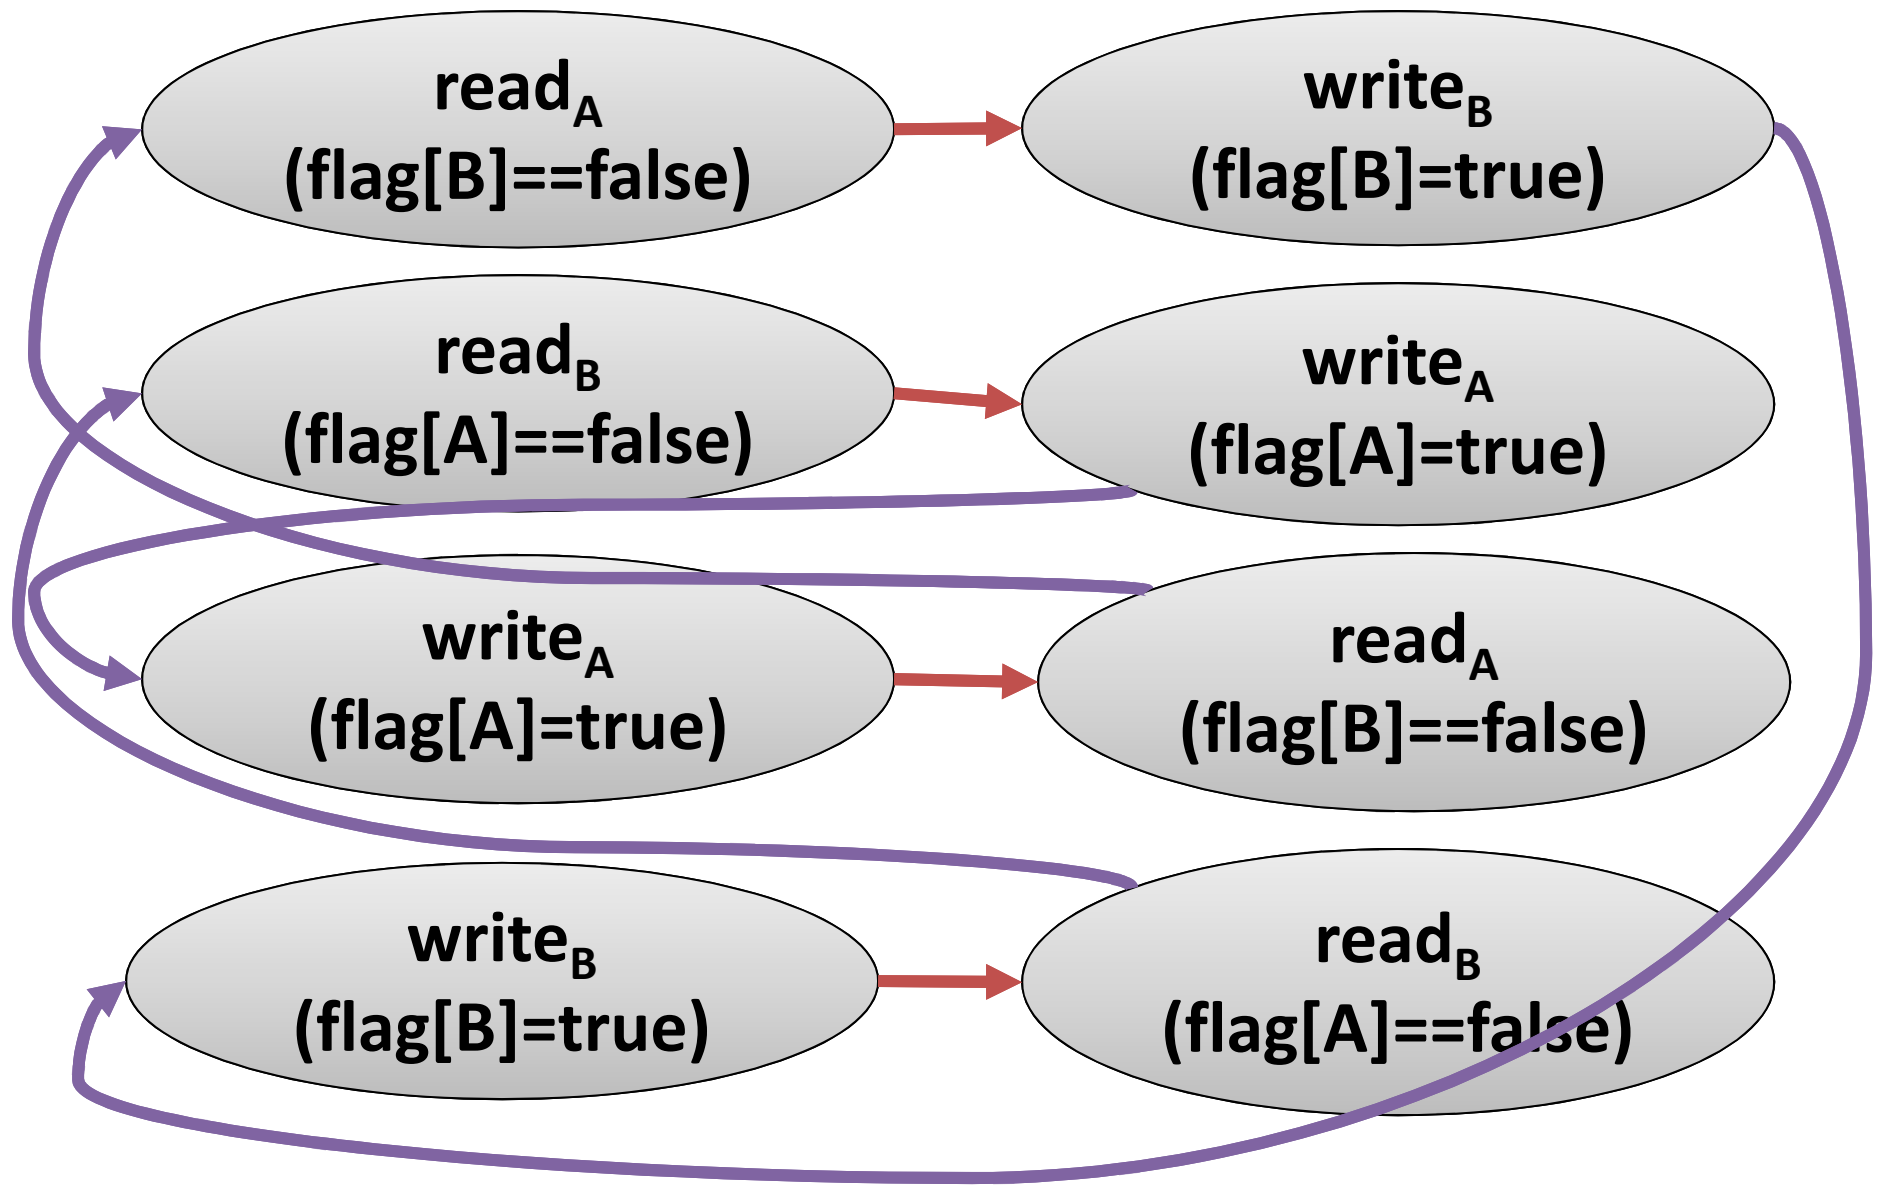
- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

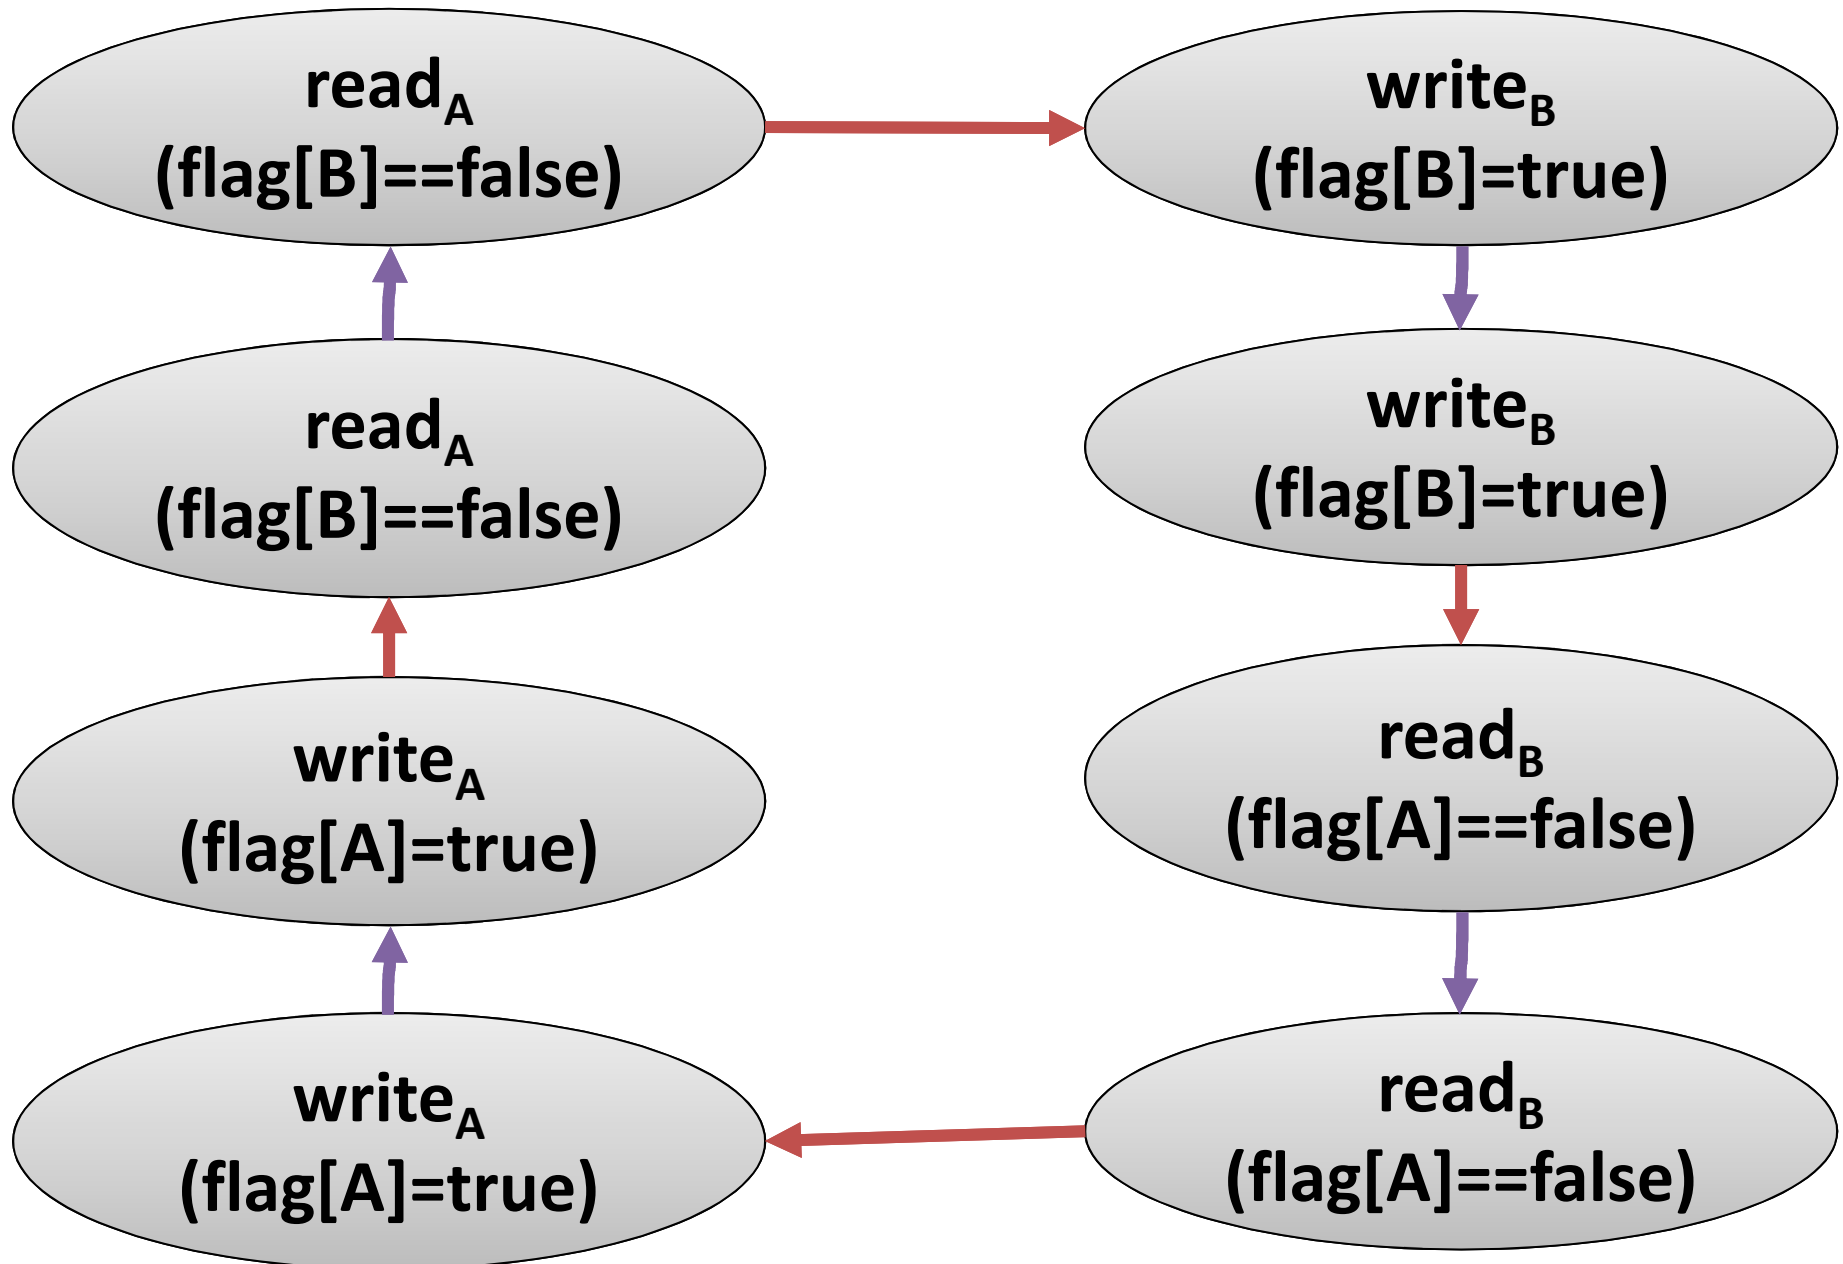
- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Combining

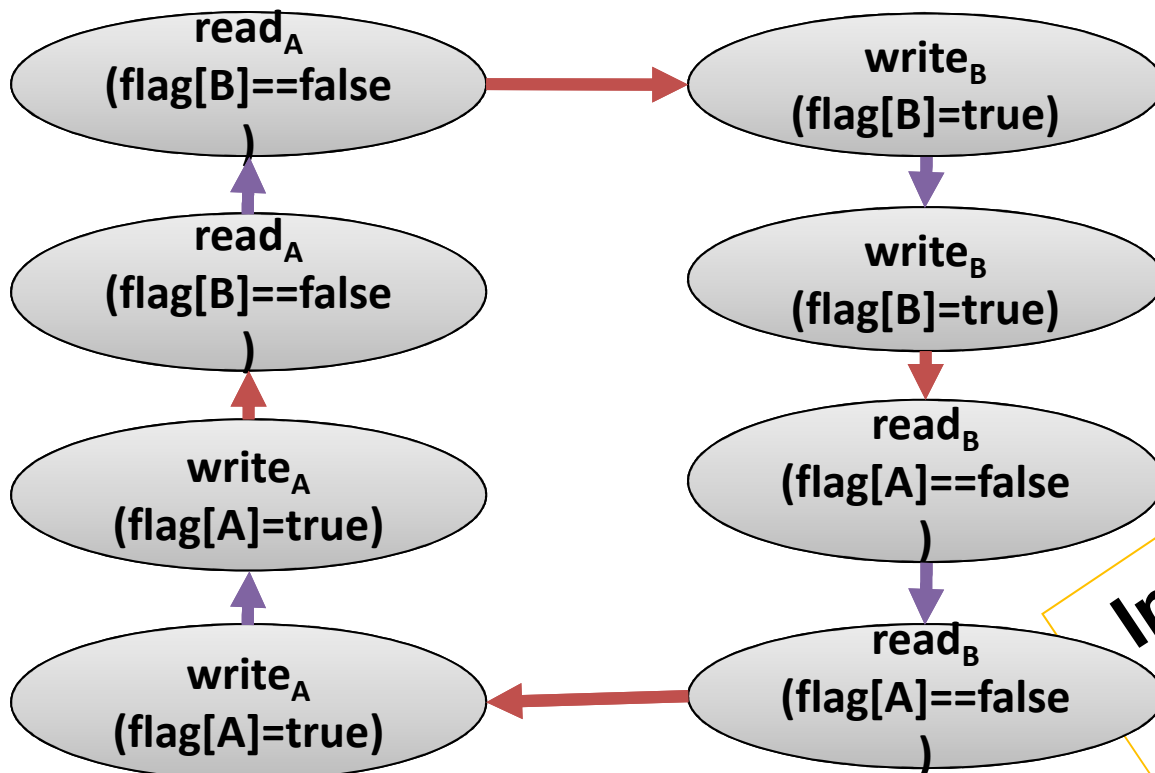


Combining: Rearranging



Cycle!

Satisfy
Mutual Execution
Proved



Impossible in a partial order

Deadlock Freedom: Progress

- LockOne Fails deadlock-freedom
 - Sequential executions OK
 - Concurrent execution

```
//Thread A
```

```
flag[i] = true;  
while (flag[j]){}
```

```
//Thread B
```

```
flag[j] = true;  
while (flag[i]){}
```

Deadlock Freedom: Progress

- LockOne Fails deadlock-freedom
 - Sequential executions OK
 - Concurrent execution can deadlock

```
//Thread A
```

```
flag[i] = true;
```

```
while (flag[j]){}
```

```
//Thread B
```

```
flag[j] = true;
```

```
while (flag[i]){}
```



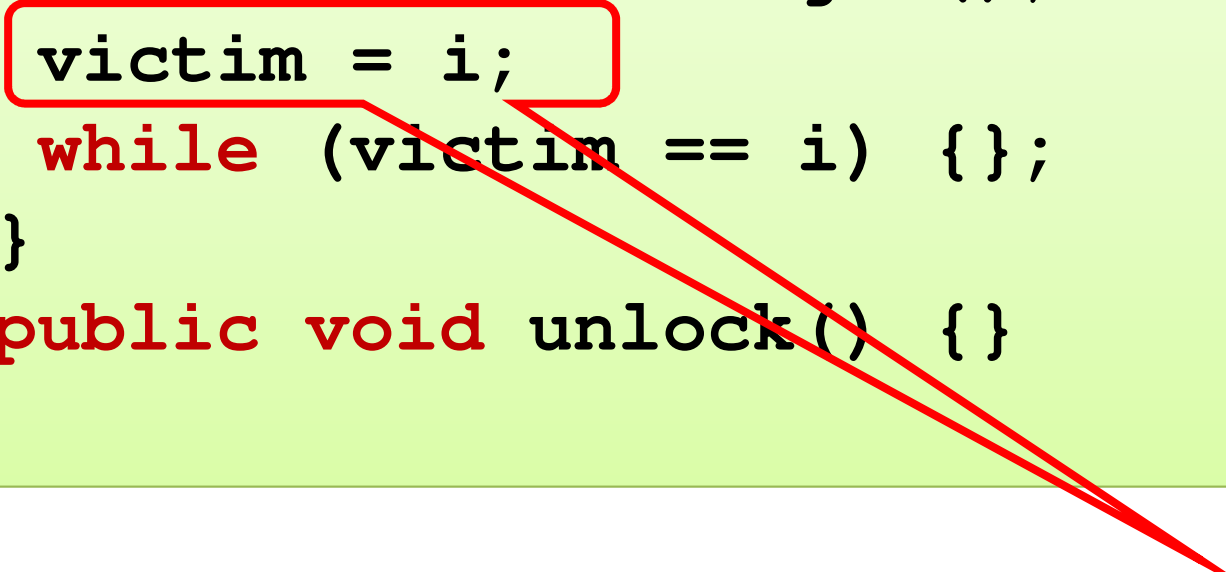
Both will wait for each other
to finish

Lock Two: Algorithm 2

```
class LockTwo {  
    int victim;  
    public void lock() {  
        int i = ThreadID.get();  
        victim = i;  
        while (victim == i) {}  
    }  
    public void unlock() {}  
}
```

Lock Two: Algorithm 2

```
class LockTwo {  
    int victim;  
    public void lock() {  
        int i = ThreadID.get();  
        victim = i;  
        while (victim == i) {}  
    }  
    public void unlock() {}  
}
```



Let other go first

Lock Two: Algorithm 2

```
class LockTwo {  
    int victim;  
    public void lock() {  
        int i = ThreadID.get();  
        victim = i;  
        while (victim == i) {};  
    }  
    public void unlock() {}  
}
```



**Wait for
permission**

Lock Two: Algorithm 2

```
class LockTwo {  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get();  
        victim = i;  
        while (victim == i) {}  
    }  
    public void unlock() {}  
}
```



Nothing to do

LockTwo Claims

- Satisfies mutual exclusion
 - If thread `i` in CS, then `victim == j`
 - Cannot be both 0 and 1
- Not deadlock free
 - Concurrent execution does not : (
 - Sequential execution deadlocks

```
public void LockTwo() {  
    victim = i;  
    while (victim == i) {}  
}
```

Peterson's Algorithm: Combine LockOne and LockTwo

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

Peterson's Algorithm: Combine LockOne and LockTwo

Announce I'm
interested

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == j) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Peterson's Algorithm: Combine LockOne and LockTwo

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm
interested

Defer to other

Peterson's Algorithm: Combine LockOne and LockTwo

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

Peterson's Algorithm: Combine LockOne and LockTwo

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

No longer interested

Wait while other interested & I'm the victim

Peterson's Lock: Lock 3

- Satisfy Mutual Exclusion
- Satisfy Deadlock Free
- Satisfy Starvation Free

—Proof

Mutual Exclusion

(1) $\text{write}_B(\text{Flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}
```

From the Code

Also from the Code

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}
```

Assumption

(3)

$\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$

Without Loss of Generality
(WLOG) assume **A** is the last
thread to write victim

Combining Observations

(1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$

(3) $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$

Combining Observations

(1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3) $\text{write}_B(\text{victim}=B) \rightarrow$

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$

Combining Observations

(1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3) $\text{write}_B(\text{victim}=B) \rightarrow$

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$

Thread A read $\text{flag}[B] == \text{true}$ and $\text{victim} == A$, so
Thread A could not have entered the CS

Deadlock Free

```
public void lock() {  
    while(flag[j] && victim==i) {};
```

- Thread blocked
 - only at **while** loop
 - only if other's flag is true
 - only if it is the **victim**
- IF (other's flag is false) then
one or the other not the victim

Starvation Free

- Thread *i* blocked only if *j* repeatedly re-enters so that `flag[j]==true` and `victim==i`
- When *j* re-enters
 - it sets `victim` to *j*.
 - So *i* gets in


```
public void lock() {  
    flag[i] = true; victim = i;  
    while(flag[j] && victim==i){};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Nthread Synchronization

Filter Algorithm for n Threads

There are $n-1$ “waiting rooms” called levels


- At each level
 - At least one enters level
 - At least one blocked if
many try
- Only one thread makes it through



Filter

```
class FilterLock {  
    int level[n]; // level[i] for thread i  
    int victim[n]; // victim[L] for level L  
    public FilterInit(int n) {  
        for(int i=1; i<n; i++)  
            level[i]=0;  
        }  
    ...  
}
```

Filter




Thread 2 at level 4

Filter

```
class FilterLock {  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k!=i level[k]>=L) &&  
                    victim[L]==i) {};  
        }  
    }  
    public void unlock() { level[i]=0; }  
}
```

Claim: Mutex

- Start at level $L=0$
- At most $n-L$ threads enter level L
- Mutual exclusion at level $L=n-1$



No Starvation

- Filter Lock satisfies properties:
 - Just like Peterson Alg at any level
 - So no one starves
- But what about fairness?
 - Threads can be overtaken by others**

Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- If A starts before B, then A enters before B?
- But what does “start” mean?
- Need to adjust definitions

Bakery Algorithm

- Similar to Bakery Shop
- Provides First-Come-First-Served
- How?
 - Take a “number”
 - Wait until lower numbers have been served
- Lexicographic order
 - $(a,i) > (b,j)$
 - If $a > b$, or $a = b$ and $i > j$

Bakery Algorithm


```
class BakeryLock {  
    bool flag[n];  
    int label[n];  
    public BakeryLockInit(int n) {  
        for(int i = 0; i < n; i++) {  
            flag[i]= false; label[i] = 0;  
        }  
    }  
    ...  
}
```

Bakery Algorithm

```
class Bakery Lock {
```

```
    bool flag[n];
```

```
    int label[n];
```



CS

Bakery Algorithm

```
class BakeryLock {  
    public void lock() {  
        flag[i]=true;  
        label[i]=max(label[0], ..., label[n-1])+1;  
        while (  $\exists k$  flag[k] &&  
                label[i] > label[k]) {}  
    }  
}
```

Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
 - $\text{flag}[A]$ is *false*, or
 - $\text{label}[A] > \text{label}[B]$

No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)

```
label[i] = max(label[0], ..., label[n-1]) + 1;
```